Outlook

Re: [WISE2026] Cannot see the reviews on my paper

From Claudio Silvestri <silvestri@unive.it>
Date Wed 26/08/2026 12:32
To Enrico.Daga <enrico.daga@open.ac.uk>
Cc Yannis Manolopoulos <manolopo@csd.auth.gr>; bhli@nuaa.edu.cn <bhli@nuaa.edu.cn>; richard.chbeir@univ-pau.fr <richard.chbeir@univ-pau.fr>; mbarhamgi@qu.edu.qa <mbarhamgi@qu.edu.qa>

Paper ID	514
Paper Title	From Fragmented Records to Actionable Intelligence: LACK, a Knowledge Graph for Climate Accountability Practitioners
Track Name	Industry Papers

Reviewer #1

Questions

1. Overall Rating
Weak Accept

2. Review confidence
Reviewer is knowledgeable

3. Detailed Comments
Overview
This paper presents LACK, a knowledge graph for climate lobbying accountability. It is built from DeSmog and LobbyMap. It contains ~28k entities and ~200k typed relations. The paper covers ontology design, an LLM extraction pipeline, quality evaluation, case studies, and an expert survey. The paper is well-motivated and technically solid. Some parts of the methodology are circular. Some key findings lack correctness checks. Some evaluations are too small.

Strengths
The motivation is clear and well-supported by prior work. The extraction and entity-linking pipeline design is technically reasonable.

Weaknesses
1. Ontology design is somewhat circular. The authors did not collect requirements directly from experts. They used the source websites as a proxy for user needs. They then used an expert survey to validate this choice. But the survey also revealed gaps outside the source-derived schema (e.g., links to the political sphere). This shows the blind spot of the original design choice. This trade-off needs clearer discussion.

2. The high NIL rate (63.5%) lacks error analysis. The paper only shows a 95.7% linking rate for well-known entities as indirect support. It does not manually check

unlinked entities. So it cannot tell true absences from pipeline misses.

3. The COP30 cross-reference uses heuristic label matching. But its accuracy is not evaluated. Two key findings depend on this matching. No precision or recall is reported. No manual check is reported either.

4. The extraction quality evaluation is small and single-source. Only 300 relations from 8 DeSmog pages were checked. The full graph has 47,450 asserted relations. LobbyMap extraction quality was not evaluated separately.

5. The entity linking evaluation is also limited. Only 225 entities were checked against a gold standard. The paper does not explain the sampling method. It is unclear if both entity types and both outcomes (linked/NIL) are covered.

6. The expert survey has a small sample and a low response rate. It was sent to 25 organizations but only 9 people responded. The authors are transparent about this. But they do not discuss possible self-selection bias or the actual response rate.

7. There is no direct comparison with concurrent work. The paper mentions the UCL AI-Powered Climate Lobbying Database. But it gives no quantitative or qualitative comparison. This weakens the "first-of-its-kind" claim.

8. The 27,757 cross-source paths are not quality-checked. Only two example cases are shown. There is no sampling analysis of how many paths are actually meaningful. Some paths may be spurious or weakly related.

9. Formatting issues. The running header shows a placeholder text: "Title Suppressed Due to Excessive Length." Figure 1 appears to have reversed or rotated text. There are typos: "continuos" should be "continuous," and "where involved" should be "were involved." The data collection date (December 2025) and citation access date (April 2026) are inconsistent. Reference [13] lists the same DOI twice.

Reviewer #2

Questions

1. Overall Rating
Weak Reject

2. Review confidence
Reviewer is an expert

3. Detailed Comments
1. Paper Significance: Below Average

This paper addresses the fragmentation of information sources and the difficulty of conducting unified relationship retrieval in the domains of climate lobbying and climate accountability. It constructs a domain-specific knowledge graph named LACK. The work has clear practical relevance, and the release of the knowledge graph, the SPARQL query endpoint, and the integration of relationships across multiple sources constitute a potentially useful resource contribution.

However, from a technical perspective, the paper mainly combines existing techniques, including large language model-based relation extraction, entity linking, RDF/OWL ontology modeling, provenance representation, and rule-based reasoning. The methodological novelty is therefore relatively limited. The main novelty lies in the specific application domain and the construction of the data resource. At present, however, the paper does not provide sufficient experimental evidence or user studies to demonstrate that the resource can substantially improve the practical efficiency of climate accountability practitioners. Therefore, I consider the work potentially valuable, but its contribution beyond the current state of the art requires stronger empirical support.

2. Paper Clarity: Above Average

The paper is generally well structured. The research motivation, system construction process, resource statistics, case studies, and expert evaluation are presented in a reasonably natural order, and most sections are easy to follow. However, several key concepts and statistical definitions are not specified rigorously enough. For example, the distinctions among “relation,” “relation triple,” “fact,” and “associatedWith connectivity” remain unclear. In addition, the claim of providing “actionable intelligence” is stronger than what is supported by the current evidence.

3. Strong Points

S1. The paper addresses a clearly defined and practically relevant problem: information concerning climate lobbying is distributed across multiple investigative platforms and is difficult to retrieve and analyze as an integrated network.

S2. LACK represents not only entities and relationships but also relationship provenance, extraction activities, and partial temporal information. These features are important for investigative research and accountability applications.

S3. The authors make the ontology, knowledge graph dump, SPARQL endpoint, and construction code publicly available, providing a useful degree of openness and potential reproducibility.

S4. DeSmog and LobbyMap provide partially complementary coverage. Integrating the two sources may reveal cross-source relationships that are difficult to identify from either source independently.

S5. In addition to reporting automatic extraction quality, the paper involves investigative journalists, fact-checkers, researchers, and NGO practitioners in assessing user needs and the potential value of the resource. This reflects an awareness of real-world user requirements.

4. Weak Points

W1. Fundamental concern: The paper lacks an end-to-end recall evaluation. The current relation extraction evaluation mainly examines whether the relations already produced by the system are correct, which approximately measures precision. It does not assess how many relations present in the original pages were missed. Therefore, the current experiment cannot demonstrate that the graph provides adequate coverage of the source data.

W2. Fundamental concern: The evaluation dataset is small and insufficiently representative. Relation extraction is evaluated using only 300 extracted relations from eight DeSmog pages. The evaluation does not cover LobbyMap and does not report results by relation type. Since pages from different sources may differ substantially in structure and writing style, the reported results are unlikely to represent the quality of the entire knowledge graph.

W3. Fundamental concern: The paper lacks validation in actual practitioner workflows. The paper claims that LACK can provide “actionable intelligence,” but it does not use realistic user tasks to demonstrate improvements in information retrieval speed, investigative accuracy, information coverage, or decision quality. The expert questionnaire mainly shows that practitioners consider this type of resource valuable; it does not demonstrate that LACK itself is already effective in practice.

W4. Fundamental concern: The evidence of system operation expected for an Industry Track paper is insufficient. The paper does not report the number of actual users, query volume, operational duration, extraction cost, system latency, update frequency, error-correction mechanisms, or integration with existing practitioner workflows. In its current form, the work is closer to a resource paper or an applied knowledge graph paper than to a sufficiently validated industrial system paper.

W5. Addressable limitation: The description of the number of relations in the abstract and main text may be misleading.
The paper emphasizes that the knowledge graph contains approximately 200,000 typed relations. However, only 47,450 are asserted relations, while most of the remaining triples are generated through inverse relations, symmetric closure, and subproperty inference. Inferred triples should not be presented as equivalent to independently extracted facts.

W6. Other Comments

- “continuos updates” should be changed to “continuous updates”;
- “Nine domain experts where involved” should be changed to “Nine domain experts were involved”;
- “valuable info” should preferably be changed to “valuable information”;
- “Knoweldge Graph” should be corrected to “Knowledge Graph”

Reviewer #3

Questions

1. Overall Rating
Accept

2. Review confidence
Reviewer is knowledgeable

3. Detailed Comments
D1. LACK addresses a clear and practically important need by integrating fragmented climate-lobbying records into a queryable knowledge graph for accountability practitioners.

D2. The released resource is substantial, containing about 28K entities and over 200K triples with provenance, temporal annotations, and links to external knowledge bases.

D3. The combination of real-world case studies, competency-question coverage, and practitioner feedback provides convincing evidence of practical usefulness.

D4. The evaluation is somewhat limited by the small expert sample of nine participants and the relatively small manually validated extraction samples.

D5. Temporal information remains a notable weakness, with temporal extraction accuracy and entity-level temporal coverage substantially below other dimensions.

D6. Overall, this is a mature, openly available, and practically valuable industry contribution with clear potential for real-world climate-accountability workflows.

Reviewer #4 (Yinan Liu)| Edit Review | Delete Review

2026-08-08 04:05:50 CEST

Questions

1. Overall Rating
Weak Accept

2. Review confidence
Reviewer is knowledgeable

3. Detailed Comments
This paper presents LACK, an openly released knowledge graph that integrates climate-lobbying information from DeSmog and LobbyMap. The resource contains approximately 28,000 entities and 200,000 typed relations, provides statement-level provenance and temporal annotations, links entities to Wikidata and DBpedia, and is available through a graph dump, SPARQL endpoint, code, and an interactive interface. The problem is important for investigative journalists, NGOs, and climate-accountability practitioners, and the cross-source case studies clearly illustrate the potential value of converting fragmented records into a queryable network. The ontology and provenance design are also thoughtful and appropriate for this high-stakes domain.

My main concern is that the current evaluation does not yet establish graph-wide reliability. Relation extraction is assessed on only 300 extracted relations from eight DeSmog pages, without evaluating LobbyMap, recall, performance by relation type, or inter-annotator agreement before discussion. Entity linking is evaluated on 225 entities, but results are not separated by persons, collectives, long-tail entities, and true NIL cases; therefore, the claim that most unlinked entities are absent from general knowledge bases rather than missed by the pipeline remains insufficiently supported. The paper also lacks matched baselines or ablations for relation extraction and entity linking, such as comparison with a standard system like ReFinED or simpler variants without LLM disambiguation and web-search fallback.

In addition, the increase from 47,450 asserted relations to 188,862 relations after inference, together with 93.5% reachability through `lack:associatedWith`, may partly reflect inverse, sub-property, and symmetric materialisation rather than independently validated coverage. Asserted and inferred graph statistics should therefore be reported separately. The COP30 case study relies on heuristic label matching without precision, recall, or a manual audit, so the conclusion that some groups do not engage with the UNFCCC track should be narrowed to absence from the examined records. The expert survey is useful but small, and the reported sample size of nine respondents is difficult to reconcile with Likert means such as 4.62, 4.75, and 4.38 unless some questions have missing responses; item-level sample sizes and response distributions should be provided. The paper also does not report construction time, hardware, LLM calls, token or monetary cost, or update cost. Because incorrect relations may cause reputational harm, the interface should provide confidence information and correction or dispute mechanisms.

Overall, LACK is a valuable and well-motivated demonstration with reuse potential while encouraging the authors to narrow the strongest claims and improve the quality evaluation, baseline comparison, cost reporting, and usability evidence in the final version.

--
Claudio Silvestri
Associate professor
Università Ca' Foscari Venezia
https://www.silv.eu/

Il giorno mer 26 ago 2026 alle ore 12:14 Microsoft CMT <noreply@msr-cmt.org> ha scritto:

> User: enrico.daga@open.ac.uk
>
> Dear chairs,
>
> I was going to prepare the camera-ready, but I can only see the decision, not the contents of the reviews.
>
> The paper is 514 - From Fragmented Records to Actionable Intelligence: LACK, a Knowledge Graph for Climate Accountability Practitioners
>
> Would it be possible to send them via email: enrico.daga@open.ac.uk
>
> Many thanks,
>
> Enrico
>
>